

Image_Super_Resolution

September 5, 2020

Image Super Resolution using Autoencoders

0.1 Project Overview and Import Libraries

```
[23]: from tensorflow.keras.layers import Input, Dense, Conv2D, MaxPooling2D, Dropout
      from tensorflow.keras.layers import Conv2DTranspose, UpSampling2D, add
      from skimage.transform import resize, rescale
      from tensorflow.keras.models import Model
      from tensorflow.keras import regularizers
      import matplotlib.pyplot as plt
      from scipy import ndimage, misc
      from matplotlib import pyplot
      import tensorflow as tf
      import numpy as np
      np.random.seed(0)
      import re
      import os

      print(tf.__version__)
```

2.3.0

0.2 What are Autoencoders?

Credit: Autoencoder Schema by Francois Chollet, 2016.

Encoder Architecture

0.3 Build the Encoder

```
[24]: input_img = Input(shape=(256, 256, 3))
```

```
[25]: l1 = Conv2D(64, (3, 3), padding = 'same', activation = 'relu',
               activity_regularizer = regularizers.l1(10e-10))(input_img)
```

```

12 = Conv2D(64, (3, 3), padding = 'same', activation = 'relu',
           activity_regularizer = regularizers.l1(10e-10))(11)

13 = MaxPooling2D(padding = 'same')(12)
13 = Dropout(0.3)(13)

14 = Conv2D(128, (3, 3), padding = 'same', activation = 'relu',
           activity_regularizer = regularizers.l1(10e-10))(13)

15 = Conv2D(128, (3, 3), padding = 'same', activation = 'relu',
           activity_regularizer = regularizers.l1(10e-10))(14)

16 = MaxPooling2D(padding = 'same')(15)

17 = Conv2D(256, (3, 3), padding = 'same', activation = 'relu',
           activity_regularizer = regularizers.l1(10e-10))(16)

encoder = Model(input_img, 17)

```

[26]: `encoder.summary()`

Model: "functional_7"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 256, 256, 3)]	0
conv2d_15 (Conv2D)	(None, 256, 256, 64)	1792
conv2d_16 (Conv2D)	(None, 256, 256, 64)	36928
max_pooling2d_4 (MaxPooling2D)	(None, 128, 128, 64)	0
dropout_2 (Dropout)	(None, 128, 128, 64)	0
conv2d_17 (Conv2D)	(None, 128, 128, 128)	73856
conv2d_18 (Conv2D)	(None, 128, 128, 128)	147584
max_pooling2d_5 (MaxPooling2D)	(None, 64, 64, 128)	0
conv2d_19 (Conv2D)	(None, 64, 64, 256)	295168
Total params: 555,328		
Trainable params: 555,328		
Non-trainable params: 0		

0.4 Build the Decoder to Complete the Network

```
[27]: 11 = Conv2D(64, (3, 3), padding = 'same', activation = 'relu',
      activity_regularizer = regularizers.l1(10e-10))(input_img)

      12 = Conv2D(64, (3, 3), padding = 'same', activation = 'relu',
      activity_regularizer = regularizers.l1(10e-10))(11)

      13 = MaxPooling2D(padding = 'same')(12)
      13 = Dropout(0.3)(13)

      14 = Conv2D(128, (3, 3), padding = 'same', activation = 'relu',
      activity_regularizer = regularizers.l1(10e-10))(13)

      15 = Conv2D(128, (3, 3), padding = 'same', activation = 'relu',
      activity_regularizer = regularizers.l1(10e-10))(14)

      16 = MaxPooling2D(padding = 'same')(15)

      17 = Conv2D(256, (3, 3), padding = 'same', activation = 'relu',
      activity_regularizer = regularizers.l1(10e-10))(16)

[28]: 18 = UpSampling2D()(17)

      19 = Conv2D(128, (3, 3), padding = 'same', activation = 'relu',
      activity_regularizer = regularizers.l1(10e-10))(18)

      110 = Conv2D(128, (3, 3), padding = 'same', activation = 'relu',
      activity_regularizer = regularizers.l1(10e-10))(19)

      111 = add([15, 110])
      112 = UpSampling2D()(111)

      113 = Conv2D(64, (3, 3), padding = 'same', activation = 'relu',
      activity_regularizer = regularizers.l1(10e-10))(112)

      114 = Conv2D(64, (3, 3), padding = 'same', activation = 'relu',
      activity_regularizer = regularizers.l1(10e-10))(113)

      115 = add([114, 12])

      decoded = Conv2D(3, (3, 3), padding = 'same',
      activation = 'relu', activity_regularizer = regularizers.
      ↪l1(10e-10))(115)
```

```
autoencoder = Model(input_img, decoded)
autoencoder_hfenn = Model(input_img, decoded)
```

```
[29]: autoencoder.summary()
```

```
Model: "functional_9"
```

```
-----
Layer (type)                 Output Shape          Param #   Connected to
=====
input_2 (InputLayer)         [(None, 256, 256, 3) 0
-----
conv2d_20 (Conv2D)           (None, 256, 256, 64) 1792      input_2[0][0]
-----
conv2d_21 (Conv2D)           (None, 256, 256, 64) 36928     conv2d_20[0][0]
-----
max_pooling2d_6 (MaxPooling2D) (None, 128, 128, 64) 0          conv2d_21[0][0]
-----
dropout_3 (Dropout)          (None, 128, 128, 64) 0
max_pooling2d_6[0][0]
-----
conv2d_22 (Conv2D)           (None, 128, 128, 128 73856     dropout_3[0][0]
-----
conv2d_23 (Conv2D)           (None, 128, 128, 128 147584     conv2d_22[0][0]
-----
max_pooling2d_7 (MaxPooling2D) (None, 64, 64, 128) 0          conv2d_23[0][0]
-----
conv2d_24 (Conv2D)           (None, 64, 64, 256) 295168     max_pooling2d_7[0][0]
-----
up_sampling2d_2 (UpSampling2D) (None, 128, 128, 256 0          conv2d_24[0][0]
-----
conv2d_25 (Conv2D)           (None, 128, 128, 128 295040     up_sampling2d_2[0][0]
-----
```

```

-----
conv2d_26 (Conv2D)          (None, 128, 128, 128) 147584      conv2d_25[0][0]
-----
add_2 (Add)                 (None, 128, 128, 128) 0          conv2d_23[0][0]
                                conv2d_26[0][0]
-----
up_sampling2d_3 (UpSampling2D) (None, 256, 256, 128) 0          add_2[0][0]
-----
conv2d_27 (Conv2D)          (None, 256, 256, 64) 73792
up_sampling2d_3[0][0]
-----
conv2d_28 (Conv2D)          (None, 256, 256, 64) 36928      conv2d_27[0][0]
-----
add_3 (Add)                 (None, 256, 256, 64) 0          conv2d_28[0][0]
                                conv2d_21[0][0]
-----
conv2d_29 (Conv2D)          (None, 256, 256, 3) 1731      add_3[0][0]
=====
Total params: 1,110,403
Trainable params: 1,110,403
Non-trainable params: 0

```

```
[30]: autoencoder.compile(optimizer = 'adadelta', loss = 'mean_squared_error')
```

0.5 Create Dataset and Specify Training Routine

```
[31]: def train_batches(just_load_dataset=False):

    batches = 256

    batch = 0
    batch_nb = 0
    max_batches = -1

    ep = 4
```

```

images = []
x_train_n = []
x_train_down = []

x_train_n2 = []
x_train_down2 = []

for root, dirnames, filenames in os.walk("data/cars_train"):
    for filename in filenames:
        if re.search("\.(jpg|jpeg|JPEG|png|bmp|tiff)$", filename):
            if batch_nb == max_batches:
                return x_train_n2, x_train_down2
            filepath = os.path.join(root, filename)
            image = pyplot.imread(filepath)
            if len(image.shape) > 2:

                image_resized = resize(image, (256, 256)) # Resize the
↪image so that every image is the same size
                x_train_n.append(image_resized) # Add this image to the
↪high res dataset
                dwn1 = rescale(image_resized, 2)
                x_train_down.append(rescale(dwn1,0.5))
                batch += 1
            if batch == batches:
                batch_nb += 1

                x_train_n2 = np.array(x_train_n)
                x_train_down2 = np.array(x_train_down)

            if just_load_dataset:
                return x_train_n2, x_train_down2

            print('Training batch', batch_nb, '(', batches, ')')

            autoencoder.fit(x_train_down2, x_train_n2,
                            epochs=ep,
                            batch_size=10,
                            shuffle=True,
                            validation_split=0.15)

            x_train_n = []
            x_train_down = []

            batch = 0

return x_train_n2, x_train_down2

```

0.6 Load the Dataset and Pre-trained Model

```
[32]: x_train_n, x_train_down = train_batches(just_load_dataset = True)
```

```
[33]: autoencoder.load_weights('data/sr.img_net.mse.final_model5.no_patch.weights.  
→best.hdf5')
```

0.7 Model Predictions and Visualizing the Results

```
[34]: encoder.load_weights('data/encoder_weights.hdf5')
```

```
[35]: encoded_imgs = encoder.predict(x_train_down)
```


```
[36]: encoded_imgs.shape
```

```
[36]: (256, 64, 64, 256)
```

```
[37]: sr1 = np.clip(autoencoder.predict(x_train_down), 0.0, 1.0)
```

```
[38]: image_index = 251
```

```
[39]: plt.figure(figsize = (20, 20))  
i = 1  
ax = plt.subplot(10, 10, i)  
plt.imshow(x_train_down[image_index])  
i += 1  
ax = plt.subplot(10, 10, i)  
plt.imshow(x_train_down[image_index], interpolation = "bicubic")  
i += 1  
ax = plt.subplot(10, 10, i)  
plt.imshow(encoded_imgs[image_index].reshape((64*64, 256)))  
i += 1  
ax = plt.subplot(10, 10, i)  
plt.imshow(sr1[image_index])  
i += 1  
ax = plt.subplot(10, 10, i)  
plt.imshow(x_train_n[image_index])  
plt.show()
```



[]: